Machine Learning Nanodegree
Capstone Report

David A. Robles

January 31, 2017

1 Definition

1.1 Project Overview

Content.

1.2 Problem Statement

In this project, we will use reinforcement learning with deep learning to make an agent learn to play the game of Connect 4[1] by playing games against itself. In other words, using the formalism used by Mitchell (1997) to define a machine learning problem:

- **Task:** Playing Connect 4.

- **Performance:** Percent of games won against other agents, and accuracy of the predictions on a Connect 4 dataset.

- **Experience:** Games played against itself.

- **Target function:** $Q^\pi : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$, where $\mathcal{S}$ is the set of *states* (board positions) and $\mathcal{A}$ is the set of *actions* (moves), and $\mathbb{R}$ represents the value of being in a state $s \in \mathcal{S}$, applying a action $a \in \mathcal{A}$, and following policy π thereafter.

- **Target function representation:** Deep neural network.

Therefore, I seek to build a Q-learning agent trained via a deep convolutional neural network to approximate the optimal action-value function:

$$Q^*(s, a) = \max_\pi Q^\pi(s, a), \forall s \in \mathcal{S}, a \in \mathcal{A} \tag{1}$$

which is the maximum sum of rewards achievable by a behaviour policy π.

1.3 Metrics

- **Winning percentage.** This metric consists in playing a high number of games (e.g. 100,000) against another agent (e.g. a random agent), and calculating the average of games won by the agent that uses the learned value function.

- **Prediction accuracy.** The learned value function will be used to predict the game-theoretic outcomes (win, loss or draw) of the board positions in the Connect 4 Data Set.

---

[1] https://en.wikipedia.org/wiki/Connect_Four

Accuracy is not a reliable metric for the real performance of a classifier, because it will yield misleading results if the data set is unbalanced (that is, when the number of samples in different classes vary greatly). For example, if there were 95 cats and only 5 dogs in the data set, the classifier could easily be biased into classifying all the samples as cats. The overall accuracy would be 95%, but in practice the classifier would have a 100% recognition rate for the cat class but a 0% recognition rate for the dog class.

2 Analysis

2.1 Data Exploration

There are two types of data used in this project: a) data generated by two games used as reinforcement learning environments used for training, and b) a dataset used for the testing phase.

2.1.1 Tic-Tac-Toe Environment

Implementation

Tic-Tac-Toe is a paper-and-pencil game for two players, X and O, who take turns marking the spaces in a 3x3 grid. The player who succeeds in placing three of their marks in a horizontal, vertical, or diagonal row wins the game. Figure 1 shows three Tic-Tac-Toe game positions.

(a) Win for X   (b) Win for O   (c) Draw

Figure 1: Examples of wins, losses and draws in Tic-Tac-Toe

As an example, we will simulate 10,000 games using a random players.

```
from capstone.game import TicTacToe
game = TicTacToe()
game.legal_moves() # [1, 2, 3, 4, 5, 6, 7, 8, 9]
print(game)
```

2.1.2 Connect 4 Environment

Implementation

Connect 4 is a two-player board game of perfect information where pieces are dropped into the columns of a vertical 6×7 grid with the goal of forming a straight line of 4 connected pieces. There are at most 7 actions per state, since placing a piece in a column is a legal action only if that column has at least one empty location. Figure 2 shows three Connect 4 game positions.

(a) Win for PLAYER 1   **(b)** Win for PLAYER 2   **(c)** Draw

Figure 2: Examples of wins, losses and draws in Connect Four

2.1.3 UCI Connect 4 Data Set

As part of the testing phase, we will use the *Connect 4 Data Set*[2] that is available from the UCI Machine Learning Repository (Hettich and Merz, 1998). This dataset has a total of 67,557 instances, representing all legal 8-ply positions in Connect 4 in which neither player has won yet, and which the next move is not forced. Each instance is described by 42 features, one for each space in the 6×7 board, and can belong to one of the classes $\{X, O, B\}$, where X is the first player, O is the second player, and B is empty. The outcome class is the game theoretical value for the first player, and can belong to one of the classes $\{WIN, LOSS, DRAW\}$. There are 44,473 wins, 16,635 losses and 6,449 draws.

## 2.2 Exploratory Visualization

2.2.1 UCI Connect 4 Data Set

As an example, here are ten instances selected randomly:

	Features																	Target
No.	a1	a2	a3	a4	a5	a6	b1	b2	...	f5	f6	g1	g2	g3	g4	g5	g6	outcome
0	b	b	b	b	b	b	b	b	...	b	b	b	b	b	b	b	b	win
1	b	b	b	b	b	b	b	b	...	b	b	b	b	b	b	b	b	win
2	b	b	b	b	b	b	o	b	...	b	b	b	b	b	b	b	b	win
3	b	b	b	b	b	b	b	b	...	b	b	b	b	b	b	b	b	win
4	o	b	b	b	b	b	b	b	...	b	b	b	b	b	b	b	b	win
...	..	..	..	..	..	..	..	..	...	..	..	..	..	..	..	..	..	...
67552	x	x	b	b	b	b	o	x	...	b	b	o	o	x	b	b	b	loss
67553	x	x	b	b	b	b	o	b	...	b	b	o	x	o	o	x	b	draw
67554	x	x	b	b	b	b	o	o	...	b	b	o	x	x	o	b	b	loss
67555	x	o	b	b	b	b	o	b	...	b	b	o	x	o	x	x	b	draw
67556	x	o	o	o	x	b	o	b	...	b	b	x	b	b	b	b	b	draw

Table 1: UCI Connect 4 Dataset. Each row represents a different board position, and each feature represents a specific cell in the board. The target is the outcome of the game for the first player, assuming perfect play.

Figure 3 shows a visual representation of the 10 instances of the data set.

---

[2]https://archive.ics.uci.edu/ml/datasets/Connect-4
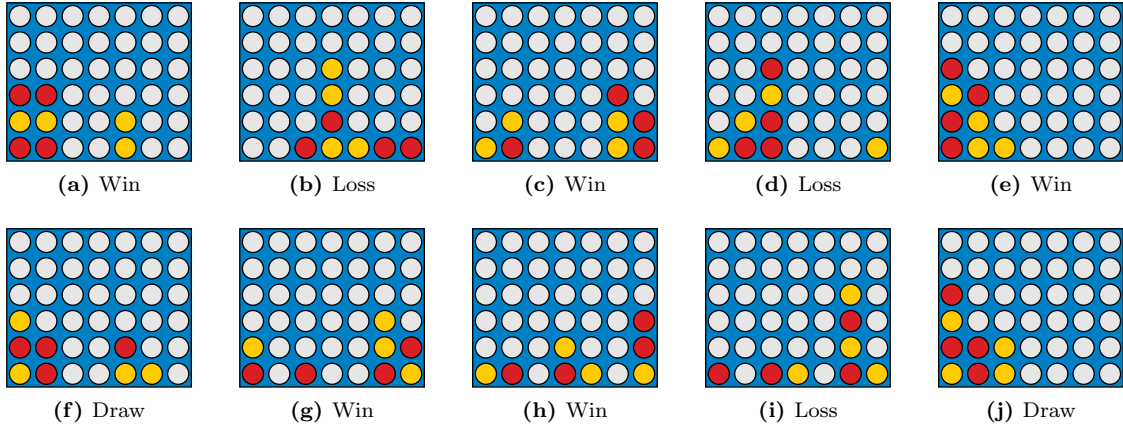
Figure 3: 10 randomly selected instances of the UCI Connect 4 Data Set. The outcome of each board position is from the point of view of the firs player (Yellow discs).

2.3 Algorithms and Techniques

2.3.1 Alpha-beta pruning

Implementation

Alpha-beta pruning (Knuth and Moore, 1975) is the most common game tree search algorithm for two-player games of perfect information. It extends the minimax algorithm to reduce the number of nodes that are evaluated in the game tree. Instead of calculating the exact minimax values for all the nodes in the game tree, alpha-beta prunes away branches that will not have any effect in the selection of the best move.

2.3.2 Q-learning

Implementation

One of the most basic and popular methods to estimate action-value functions is the *Q-learning* algorithm (Watkins, 1989). It is model-free online off-policy algorithm, whose main strength is that it is able to compare the expected utility of the available actions without requiring a model of the environment. Q-learning works by learning an action-value function that gives the expected utility of taking a given action in a given state and following a fixed policy thereafter. The update rule uses action-values and a built-in max-operator over the action-values of the next state in order to update $Q(s_t, a_t)$ as follows,

$$Q(s_t, a_t) \leftarrow Q(s_t, a_t) + \alpha[r_{t+1} + \gamma \max_a Q(s_{t+1}, a) - Q(s_t, a_t)] \tag{2}$$

The agent makes a step in the environment from state s_t to s_{t+1} using action a_t while receiving reward r_t. The update takes place on the action-value a_t in the state s_t from which this action was executed.

Q-learning is exploration-intensive, which means that it will converge to the optimal policy regardless of the exploration policy being followed, under the assumption that each state-action pair is visited an infinite number of times, and the learning parameter α is decreased appropriately (Watkins and Peter, 1992).

2.3.3 Self-play

Self-play is by far the most popular training method. It is a single policy $\pi(s, a)$ that is used by both players in a two-player game, $\pi_1(s, a) = \pi_2(s, a) = \pi(s, a)$. The first reason for its popularity is that training is quickest if the learner's opponent is roughly equally strong, and that definitely holds for self-play. As a

second reason for popularity, there is no need to implement or access a different agent with roughly equal playing strength. However, self-play has several drawbacks, with the main one being that a single opponent does not provide sufficient exploration (Szita, 2011).

2.4   Benchmark

- **Random agent**. This benchmark consists in playing against an agent that takes uniformly random moves. This is the most basic benchmark, but first we have to be sure that our learned evaluation function can play better than a random agent before moving into a harder benchmark. Also, this will help us to detect bugs in the code and algorithms: if a learned value function does not play significantly better than a random agent, is not learning. The idea is to test against this benchmark using Alpha-beta pruning at 1, 2 and 4-ply search.

- **Connect 4 Data Set**. This dataset will be used as the main benchmark. The learned value function will be used to predict the game-theoretic outcomes (win, loss or draw) for the first player in the 67,557 instances of the dataset.

3   Implementation

3.1   Markov Decision Process

Implementation

A *Markov decision process* (MDP) consist of a set of states, a set of actions, a transition function and a reward function.

- $\mathcal{S}$ is the set of *states* (state space).

- $\mathcal{A}$ is the set of *actions* (action space). The set of actions that are available in some particular state $s_t \in \mathcal{S}$ is denoted $\mathcal{A}(s_t)$, such that $\mathcal{A}(s_t) \in \mathcal{P}(\mathcal{A})$, where $\mathcal{P}(\cdot)$ denotes the power set.

- $T : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \to \mathbb{R}$ is the *transition function*, which is the probability given we are in state $s_t \in \mathcal{S}$, take action $a_t \in \mathcal{A}(s_t)$ and we will transition to state $s_{t+1} \in \mathcal{S}$.

- $R : \mathcal{S} \times \mathcal{A} \times \mathcal{S} \to \mathbb{R}$ is the *reward function*, which is the immediate reward received when in state $s_t \in \mathcal{S}$ action $a_t \in \mathcal{A}$ is taken and the MDP transitions to state $s_{t+1} \in \mathcal{S}$. However, it is also possible to define it either as $R : \mathcal{S} \times \mathcal{A} \to \mathbb{R}$ or $R : \mathcal{S} \to \mathbb{R}$. The first one gives rewards for performing an action a_t in a particular state s_t, and the second gives rewards when transitioning to state s_{t+1}.

3.2   Environment

Implementation

An agent does not have access to the dynamics (reward and transition function) of the MDP. However, it interacts with an *environment* by way of three signals: a *state*, which describes the state of the environment, an *action*, which allows the agent to have some impact on the environment, and a *reward*, which provides the agent with feedback on its immediate performance.

3.3   Policy

Implementation

In an MDP, the agent acts according to a policy π, which maps each state $s \in \mathcal{S}$ to an action $a \in \mathcal{A}(s)$. A policy that specifies a unique action to be performed is called a *deterministic* policy, and is defined as $\pi : \mathcal{S} \to \mathcal{A}$. On the other hand, a *stochastic* policy $\pi : \mathcal{S} \times \mathcal{A} \to [0, 1]$ selects actions according to a probability distribution, such that for each state $s \in \mathcal{S}$, it holds that $\pi(s, a) \geq 0$ and $\sum_{a \in \mathcal{A}(s)} \pi(s, a) = 1$.

The interaction between the policy used by the agent and the environment works as follows. First, it starts at an *initial state* s_0. Then, the policy π selects an action $a_0 = \pi(s_0)$ from the set of available actions $\mathcal{A}(s_0)$, and the action is executed. The environment transitions to a new state s_1 based on the transition function T with probability $T(s_0, a_0, s_1)$, and a reward $r_0 = R(s_0, a_0, s_1)$ is received. This process continues, producing a *trajectory* of experience $s_0, a_0, s_1, r_1, a_1, s_2, r_2, a_2, \ldots$. If the task is *episodic*, the process ends in a *terminal state* s_T and is restarted in the initial state. If the task is *continiuing*, the sequence of states can be extended indefinitely.

3.3.1 Random policy

Implementation

Selects:

$$\pi_{\mathrm{rand}}(s) = \operatorname*{argmax}_{a \in \mathcal{A}(s_t)} Q(s, a) \tag{3}$$

3.3.2 Greedy policy

Implementation

Implementation

The greedy tree policy selects the *max action*, which is the greedy action with the highest value.

$$\pi_{\mathrm{greedy}}(s) = \operatorname*{argmax}_{a \in \mathcal{A}(s_t)} Q(s, a) \tag{4}$$

3.3.3 ϵ-greedy

Implementation

The ϵ-greedy tree policy selects the best action for a proportion $1 - \epsilon$ of the trials, and another action is randomly selected (with uniform probability) for a proportion ϵ,

$$\pi_\epsilon(s) = \begin{cases} \pi_{\mathrm{rand}}(s, a) & \text{if } rand() < \epsilon \\ \pi_{\mathrm{greedy}}(s, a) & \text{otherwise} \end{cases} \tag{5}$$

where $\epsilon \in [0, 1]$ and $rand()$ returns a random number from a uniform distribution $\in [0, 1]$.

3.4 Game

Informally speaking, a perfect-information game in extensive form is a tree in the sense of graph theory, in which each node represents the current state of the game, each edge represents a possible action, and the leaves represent the terminal states, which contains final outcomes through a utility function. In the field of AI these are known simply as *game trees*.

A **perfect-information extensive-form game** is defined as a tuple $G = (n, \mathcal{S}, \mathcal{S_T}, \mathcal{A}, \rho, T, U)$, where:

- n is the *number of players*;
- $\mathcal{S}$ is the set of *states*;

- $\mathcal{S}_T \subset \mathcal{S}$ is the set of *terminal states*;

- $\mathcal{A}$ is the set of *actions*. The set of actions that are available in some particular state $s_t \in \mathcal{S}$ for the player in turn $\rho(s_t)$ is denoted $\mathcal{A}(s_t)$, such that $\mathcal{A}(s_t) \in \mathcal{P}(\mathcal{A})$, where $\mathcal{P}(\cdot)$ denotes the power set.

- $\rho : \overline{\mathcal{S}_T} \to \{i \in \mathbb{N} : 1 \leq i \leq n\}$ is the *player function*, which assigns to each non-terminal state $s_t \in \overline{\mathcal{S}_T}$ a player i who chooses an action at that state;

- $T : \overline{\mathcal{S}_T} \times \mathcal{A} \to \mathcal{S}$ is the *transition function*, which maps a non-terminal state and an action to a new state;

- $U : S_T \to \mathbb{R}^n$ is the *utility function*, which returns a vector of payoffs $U(s_T) = (u_1, \ldots, u_n)$ for each of the n players of the game. We will use the notation $U_i(s_T)$ to represent the payoff for the ith player when the game is over at state $s_T \in S_T$.

In two-player perfect-information extensive-form games, a *ply* refers to one turn taken by one of the players.

4 Methodology

4.1 Data Preprocessing

Content.

4.2 Implementation

- Two games were implemented: Tic Tac Toe and Connect 4.

- Converting a to an MDP by using a fixed agent.

- Show that q-learning learns to play against a fixed opponent using Tic-Tac-Toe.

- Show that q-learning learns to play against a fixed opponent using Connect 4.

4.3 Refinement

Content.

4.4 Games

4.4.1 Tic Tac Toe

Implementation

### 4.4.2 GameMDP converter

Converting an MDP

5   Examples

5.1   Estimate state-action values using tabular Q-learning

Implementation

First, we will test our q-learning algorithm by learning the state action values for an agent that plays against a fixed tic tac toe player. this fixed tic tac toe player will be an AlphaBeta player, which means it will always make an optimal move. An optimal move means that is the best possible move in that situation.

We created a custom board position that is small but that it can demonstrate the idea of learning the q-learning values.

The situation consists where we are in the following position:

Figure 4: Tic-Tac-Toe board, s, with five legal moves, $\mathcal{A}(s) = \{1, 2, 4, 6, 9\}$.

In this position the player X is next, and it has 5 available moves: 1, 3, 6, 9.
We ran the Q-learning algorithm with the following parameters and these are the results:

(a) $Q(s, 1) = -1.0$ **(b)** $Q(s, 2) = -1.0$ **(c)** $Q(s, 4) = -1.0$ **(d)** $Q(s, 6) = -1.0$ **(e)** $Q(s, 9) = 1.0$

Figure 5

As we can see we were able to learn the state-action values for each board position. It is clear that the best board is the one that leads to a direct win, and the other boards lead to either draws (like in...) or a loss.

Now we will run the same algorithm but starting from the initial position.

As we can see it learned to play perfectly against an alpha-beta player.

With this example we can confirm that the implementation of our q-learning algorithm is working fine, and now we can move into our next step.

## 5.2   Learn Connect 4 state-action values using Q-learning
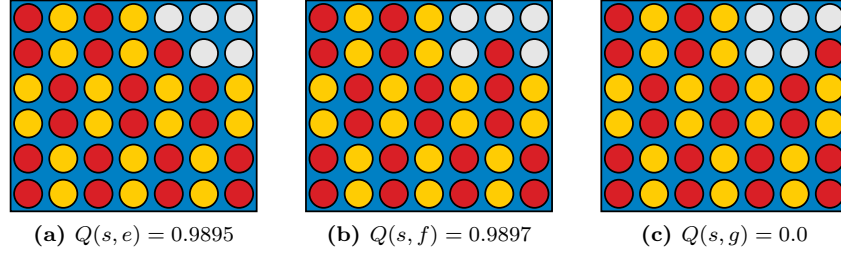
Implementation

asd fas dfasd fas dfasd

Figure 6: The predicted state-action value during training.

Figure 7: Connect 4 board, s, with three legal moves: $\{e, f, g\}$.

(a) $Q(s, e) = 0.9895$ **(b)** $Q(s, f) = 0.9897$ **(c)** $Q(s, g) = 0.0$

Figure 8: All the next states and their q values.

6 Results

6.1 Model Evaluation and Validation

Content.

6.2 Justification

Content.

7   Conclusion

7.1   Free-Form Visualization

Content.

7.2   Reflection

Content.

7.3   Improvement

Content.

References

C. B. S. Hettich and C. Merz. UCI repository of machine learning databases, 1998.

Donald E. Knuth and Ronald W. Moore. An analysis of alpha-beta pruning. *Artificial Intelligence*, 6(4), 1975.

T. M. Mitchell. *Machine Learning*. McGraw Hill, 1997.

I. Szita. Reinforcement learning and markov decision processes. In *Reinforcement Learning: State of the Art*. Springer, 2011.

Christopher J.C.H. Watkins and Dayan Peter. Q-learning. *Machine Learning*, 8:279–292, 1992.

C.J.C.H. Watkins. *Learning from Delayed Rewards*. PhD thesis, Cambridge University, 1989.